

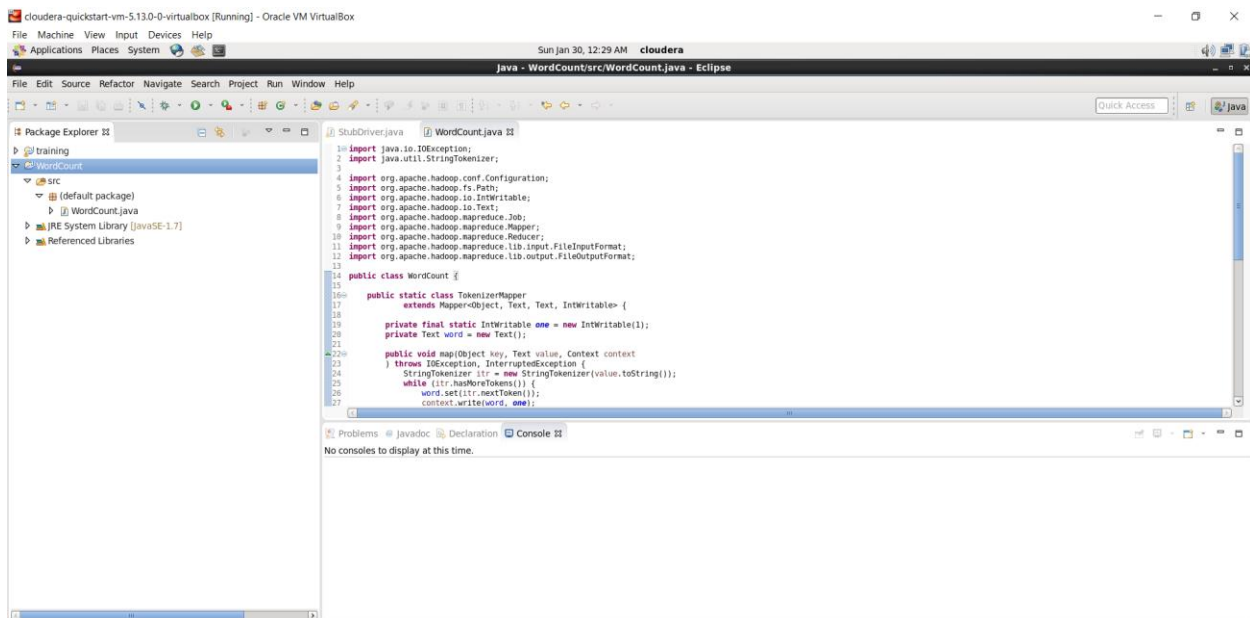
# CSCE 5300 INTRODUCTION TO BIG DATA AND DATA SCIENCE

**NAME :** Eswara Reddy Thimmapuram

## Use case Mapper Reduce word count process

- Firstly, the input is given in a specified size and then the input is split into different parts. Then the next job is to convert input block into string. So the mapper takes starting address and ending address as input.
- Then the mapper converts input into string and tokenizes into words. The mapper will then append 1 to each word for first mapper output and so on. The output of mapper would be key-value pairs. After mapper process the next jobs are sorting and shuffling.
- In next phase, all words come to their respective places in sorting phase and in shuffling phase all similar words go to similar reducer. Then the reducer will sum up values by using keys. The output of reducer will be the final result which is the sum of words that are present in the input file.

## WordCount Program



- Firstly we create new java project WordCount and import all header files and will add all required external jar files. Then we create new class named WordCount .

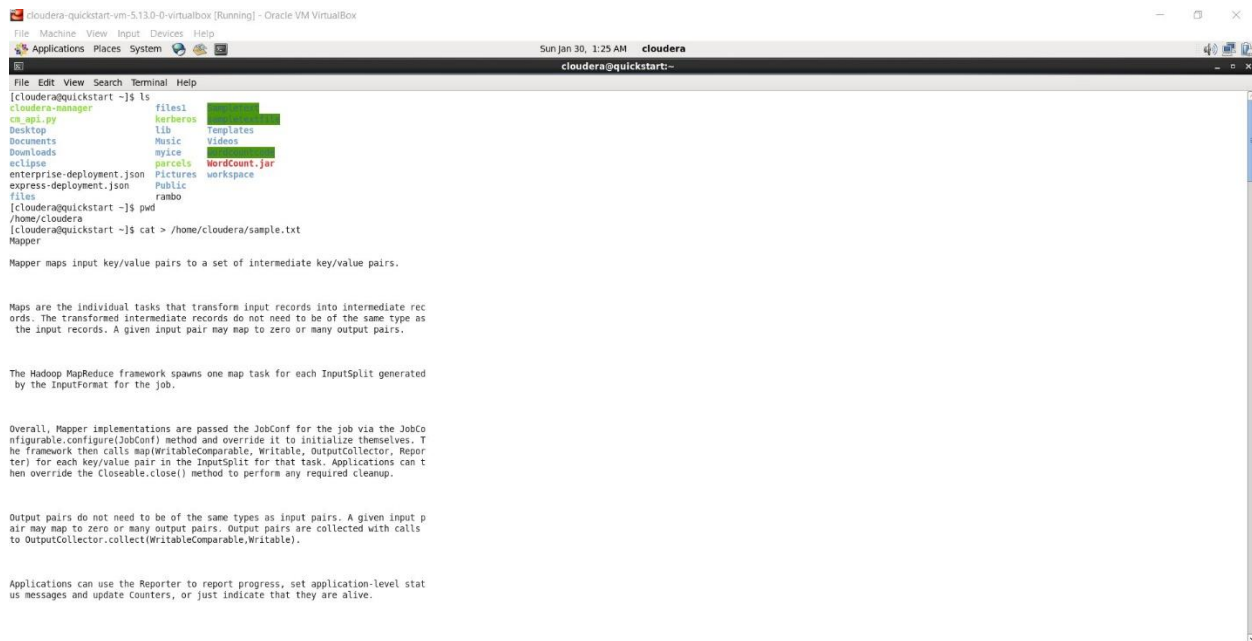
- Then we extend the mapper class by creating TokenizerMapper and it takes arguments like key-in,value-in,key-out and value-out. Then we create word variable to store word and one variable to store value.
- Then we create map function which takes arguments like key-in as key, textual content as value variable and context to emit result to next stage.
- Then we StringTokenizer which divides string into tokens and store each of them in itr variable.
- Then we define reducer class as IntSumReducer which takes four inputs. Then we create result variable to store how many number of times the key is appearing.
- Then we write reduce function which takes word as key and value as value and if the word repeats then it will be added to the sum value.
- Finally the reducer gives the final output of the program which is the total sum of words that are present in input file.

```

cloudera-quickstart-vm-5.13.0-6-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
Sun Jan 30, 1:16 AM cloudera
cloudera@quickstart:/home/cloudera/Sampletext
File Edit View Search Terminal Help
cloudera@quickstart ~$ mkdir Sampletext
cloudera@quickstart ~$ cd Sampletext
cloudera@quickstart Sampletext$ su
Password:
root@quickstart Sampletext$ mount -t vboxsf sampletextfile /home/cloudera/Sampl
text
root@quickstart Sampletext$

```

- Then we create Sampletext file to store the input we have to give to the program.



The screenshot shows a terminal window titled "cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox". The terminal prompt is "cloudera@quickstart:~". The user has run "ls" and "pwd", showing they are in the "/home/cloudera" directory. They then run "cat > /home/cloudera/sample.txt". The terminal output shows the file system structure and the result of the cat command, which is a block of text about the Mapper class in Hadoop.

```
cloudera@quickstart ~$ ls
cloudera-manager  files1  herberas
ca_api.py         lib     Music
Desktop           myrice parcels  WordCount.jar
Documents         enterprise-deployment.json  Pictures  workspace
Downloads         express-deployment.json    Public
eclipse           files  rambo
enterprise-deployment.json
express-deployment.json
files
rambo

cloudera@quickstart ~$ pwd
/home/cloudera
cloudera@quickstart ~$ cat > /home/cloudera/sample.txt
Mapper

Mapper maps input key/value pairs to a set of intermediate key/value pairs.

Maps are the individual tasks that transform input records into intermediate records. The transformed intermediate records do not need to be of the same type as the input records. A given input pair may map to zero or many output pairs.

The Hadoop MapReduce framework spawns one map task for each InputSplit generated by the InputFormat for the job.

Overall, Mapper implementations are passed the JobConf for the job via the JobConf.setConf() method and override it to initialize themselves. The framework then calls map(WritableComparable, Writable, OutputCollector, Reporter) for each key/value pair in the InputSplit for that task. Applications can then override the Closeable.close() method to perform any required cleanup.

Output pairs do not need to be of the same types as input pairs. A given input pair may map to zero or many output pairs. Output pairs are collected with calls to OutputCollector.collect(WritableComparable, Writable).

Applications can use the Reporter to report progress, set application-level status messages and update Counters, or just indicate that they are alive.
```

- Then we use cat command to store the input into sampletext file.



The screenshot shows a terminal window titled "cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox". The terminal prompt is "cloudera@quickstart:~". The user has run "ls" and "pwd", showing they are in the "/home/cloudera" directory. They then run "cat > /home/cloudera/sample.txt". The terminal output shows the file system structure and the result of the cat command, which is a block of text about the Mapper class in Hadoop.

```
cloudera@quickstart ~$ ls
cloudera-manager  files1  herberas
ca_api.py         lib     Music
Desktop           myrice parcels  WordCount.jar
Documents         enterprise-deployment.json  Pictures  workspace
Downloads         express-deployment.json    Public
eclipse           files  rambo
enterprise-deployment.json
express-deployment.json
files
rambo

cloudera@quickstart ~$ pwd
/home/cloudera
cloudera@quickstart ~$ cat > /home/cloudera/sample.txt
Mapper

Mapper maps input key/value pairs to a set of intermediate key/value pairs.

Maps are the individual tasks that transform input records into intermediate records. The transformed intermediate records do not need to be of the same type as the input records. A given input pair may map to zero or many output pairs.

The Hadoop MapReduce framework spawns one map task for each InputSplit generated by the InputFormat for the job.

Overall, Mapper implementations are passed the JobConf for the job via the JobConf.setConf() method and override it to initialize themselves. The framework then calls map(WritableComparable, Writable, OutputCollector, Reporter) for each key/value pair in the InputSplit for that task. Applications can then override the Closeable.close() method to perform any required cleanup.

Output pairs do not need to be of the same types as input pairs. A given input pair may map to zero or many output pairs. Output pairs are collected with calls to OutputCollector.collect(WritableComparable, Writable).

Applications can use the Reporter to report progress, set application-level status messages and update Counters, or just indicate that they are alive.
```

- Then we create input1 file and use put command to copy file from local system to Hadoop system.

```

Applications can use the Reporter to report progress, set application-level status messages and update Counters, or just indicate that they are alive.

The output of the Reducer is not sorted.
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/WordCount.jar WordCount input/sample.txt output1
22/01/30 01:33:21 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/01/30 01:33:24 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
22/01/30 01:33:25 INFO Input.FileInputFormat: Total input paths to process : 1
22/01/30 01:33:25 INFO mapreduce.JobSubmitter: number of splits:1
22/01/30 01:33:27 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1643520845524_0001
22/01/30 01:33:31 INFO impl.YarnClientImpl: Submitted application application_1643520845524_0001
22/01/30 01:33:32 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1643520845524_0001/
22/01/30 01:33:32 INFO mapreduce.Job: Running job: job_1643520845524_0001
22/01/30 01:34:11 INFO mapreduce.Job: Job job_1643520845524_0001 running in uber mode : false
22/01/30 01:34:59 INFO mapreduce.Job: map 100% reduce 0%
22/01/30 01:39:55 INFO mapreduce.Job: map 100% reduce 100%
22/01/30 01:39:55 INFO mapreduce.Job: Job job_1643520845524_0001 completed successfully
22/01/30 01:39:55 INFO mapreduce.Job: Counters: 49
File System Counters
  FILE: Number of bytes read=4204
  FILE: Number of bytes written=295149
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=4924
  HDFS: Number of bytes written=3035
  HDFS: Number of read operations=0
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=43587
  Total time spent by all reduces in occupied slots (ms)=107215
  Total time spent by all map tasks (ms)=43587
  Total time spent by all reduce tasks (ms)=107215
  Total vcore-milliseconds taken by all map tasks=43587
  Total vcore-milliseconds taken by all reduce tasks=107215
  Total megabyte-milliseconds taken by all map tasks=44633008
  Total megabyte-milliseconds taken by all reduce tasks=109788160
Map-Reduce Framework
  Map input records=111
  Map output records=719

```

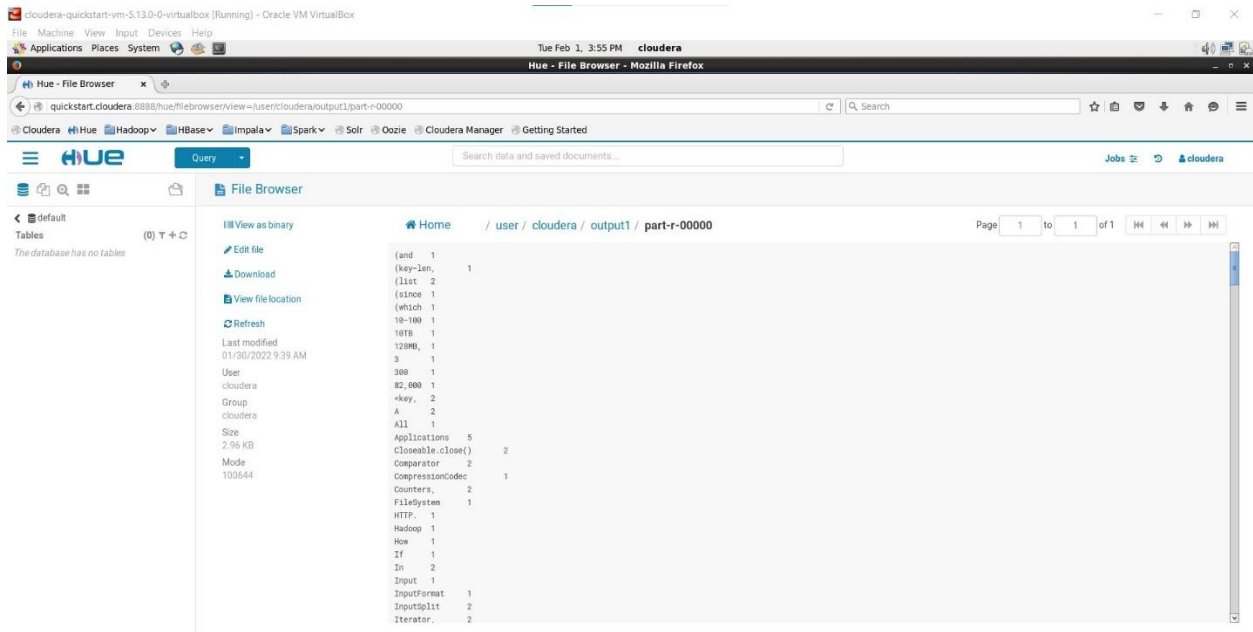
- Then we run the mapreduce job command and then the execution continues.

```

[cloudera@quickstart ~]$ hadoop fs -ls output1
Found 2 items
-rw-r--r-- 1 cloudera cloudera 0 2022-01-30 01:39 output1/SUCCESS
-rw-r--r-- 1 cloudera cloudera 3035 2022-01-30 01:39 output1/part-r-000000
[cloudera@quickstart ~]$ hadoop fs -cat output1/part-r-000000
(
  1
  (key-len, 1
  (list 2
  (since 1
  (which 1
  10-100 1
  10TB 1
  128MB, 1
  3 1
  300 1
  82,000 1
  <key, 2
  A 2
  All 1
  Applications 5
  Closeable.close() 2
  Comparator 2
  CompressionCodec 1
  Counters, 2
  FileSystem 1
  HTTP 1
  Hadoop 1
  How 1
  If 1
  In 2
  Input 1
  InputFormat 1
  InputSplit 2
  Iterator, 2
  JobConf 2
  JobConf, 1
  JobConf.setCombinerClass(Class), 1
  JobConf.setNumReduceTasks(int), 1
  JobConf.setOutputKeyComparatorClass(Class), 1
  JobConf.setOutputKeyComparatorClass(Class), 1
  JobConf.setOutputValueGroupingComparator(Class), 1
  JobConfigurable.configure(JobConf) 2
  Many 1
  MapReduce 1

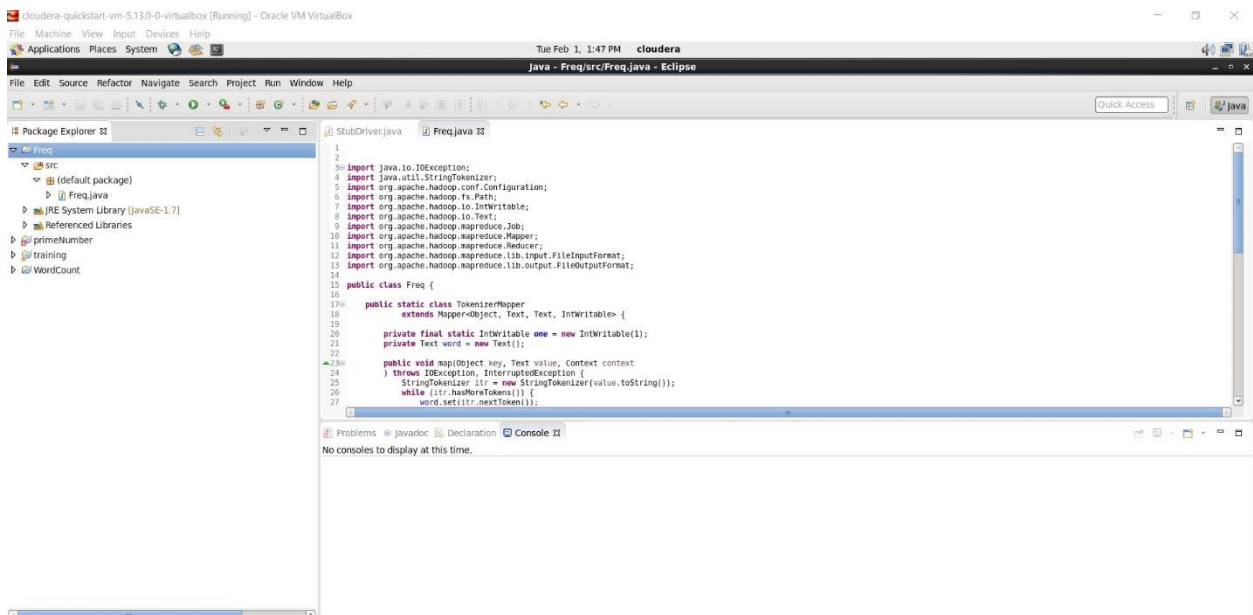
```

- Finally, we list the files in the output and use cat command to display the output which is the number of times the words are repeated in given file.



- Here is the visualized version of wordcount program

## Frequency of words with letter 'b'



- Firstly we create a java project named Freq and then import all the required header files and add all external jar libs to the program.

- Then we extend the mapper class by creating TokenizerMapper and it takes arguments like key-in,value-in,key-out and value-out. Then we create word variable to store word and one variable to store value.
- Then we create map function which takes arguments like key-in as key, textual content as value variable and context to emit result to next stage.
- Then we StringTokenizer which divides string into tokens and store each of them in its variable.
- Then we define reducer class as IntSumReducer which takes four inputs. Then we create result variable to store how many number of times the key is appearing.
- Then we write logic which adds words which starts with letter b and gives output to the reducer.

```

cloudera-quickstart-vm-5130-0-virtualbox [Running] - Oracle VM VirtualBox
File Machines View Input Devices Help
Applications Places System
Tue Feb 1, 2:04 PM cloudera
cloudera@quickstart:~$ cat > /home/cloudera/freqInput.txt
base
base
boy
dog
line
cat
[cloudera@quickstart ~]$ hadoop fs -mkdir inputFreq
[cloudera@quickstart ~]$ hadoop fs -put /home/cloudera/freqInput.txt inputFreq/
[cloudera@quickstart ~]$ hadoop fs -ls inputFreq
Found 1 items
-rw-r--r-- 1 cloudera cloudera 27 2022-02-01 13:58 inputFreq/freqInput
txt
[cloudera@quickstart ~]$ hadoop fs -cat inputFreq/freqInput.txt
base
base
boy
dog
line
cat
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Freq.jar Freq inputFreq/freqI
nput.txt output

22/02/01 14:01:22 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0
8002
22/02/01 14:01:23 WARN mapreduce.JobResourceUploader: Hadoop command-line option
parsing not performed. Implement the Tool interface and execute your applicatio
n with ToolRunner to remedy this.
22/02/01 14:01:24 INFO input.FileInputFormat: Total input paths to process : 1
22/02/01 14:01:24 INFO mapreduce.JobSubmitter: number of splits:1
22/02/01 14:01:24 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_16
43738845564_0001
22/02/01 14:01:25 INFO impl.YarnClientImpl: Submitted application application_16
43738845564_0001
22/02/01 14:01:26 INFO mapreduce.Job: The url to track the job: http://quickstar
t.cloudera:8088/proxy/application_1643738845564_0001/
22/02/01 14:01:26 INFO mapreduce.Job: Running job: job_1643738845564_0001
22/02/01 14:01:40 INFO mapreduce.Job: Job job_1643738845564_0001 running in uber mode : false
22/02/01 14:01:40 INFO mapreduce.Job: map 0% reduce 0%
22/02/01 14:01:51 INFO mapreduce.Job: map 100% reduce 0%
22/02/01 14:02:01 INFO mapreduce.Job: map 100% reduce 100%
22/02/01 14:02:02 INFO mapreduce.Job: Job job_1643738845564_0001 completed successfully
22/02/01 14:02:02 INFO mapreduce.Job: Counters: 49
File System Counters
FILE: Number of bytes read:37
FILE: Number of bytes written:286795
FILE: Number of read operations:0

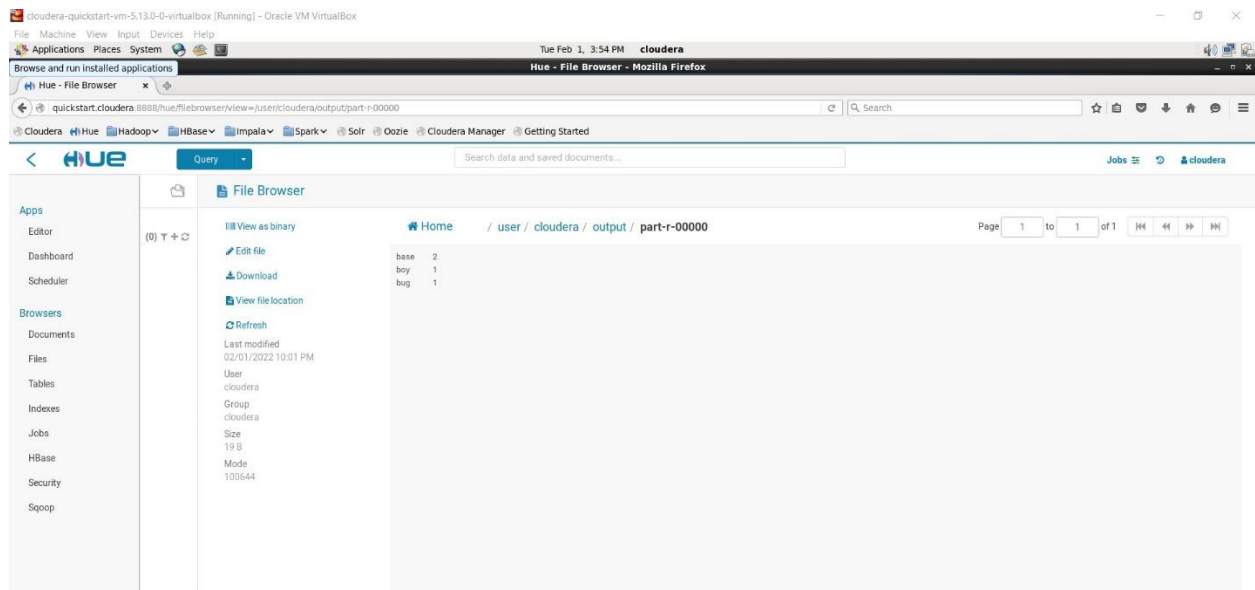
```

- Then we create a file and add required input to it using cat command. Then after we import files from local to Hadoop using put command.
- Then we display input using cat command. Then we execute map reduce job and execution continues.



```
cloudera-quickstart-vm-5.13.0-0-virtualbox (Running) - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System
Tue Feb 1, 2:05 PM cloudera
cloudera@quickstart:~$
File Edit View Search Terminal Help
Map Input records=6
Map output records=4
Map output bytes=34
Map output materialized bytes=37
Input split bytes=134
Combine input records=4
Combine output records=3
Reduce input groups=3
Reduce shuffle bytes=37
Reduce input records=3
Reduce output records=3
Spilled Records=6
Shuffled Haps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=163
CPU time spent (ms)=2340
Physical memory (bytes) snapshot=392896512
Virtual memory (bytes) snapshot=3015629544
Total committed heap usage (bytes)=226365440
Shuffle Errors
BAD_IP=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=27
File Output Format Counters
Bytes Written=19
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$ hadoop fs -ls output
Found 2 items
-rw-r--r-- 1 cloudera cloudera 0 2022-02-01 14:02 output/_SUCCESS
-rw-r--r-- 1 cloudera cloudera 19 2022-02-01 14:01 output/part-r-00000
[cloudera@quickstart ~]$ hadoop fs -cat output/part-r-00000
base 2
boy 1
bug 1
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
[cloudera@quickstart ~]$
```

- Finally, we list the files in the output and use cat command to display the output which is the number of times the words with letter ‘b’ are repeated in given file.



- Here is visualized version of frequency of b words file

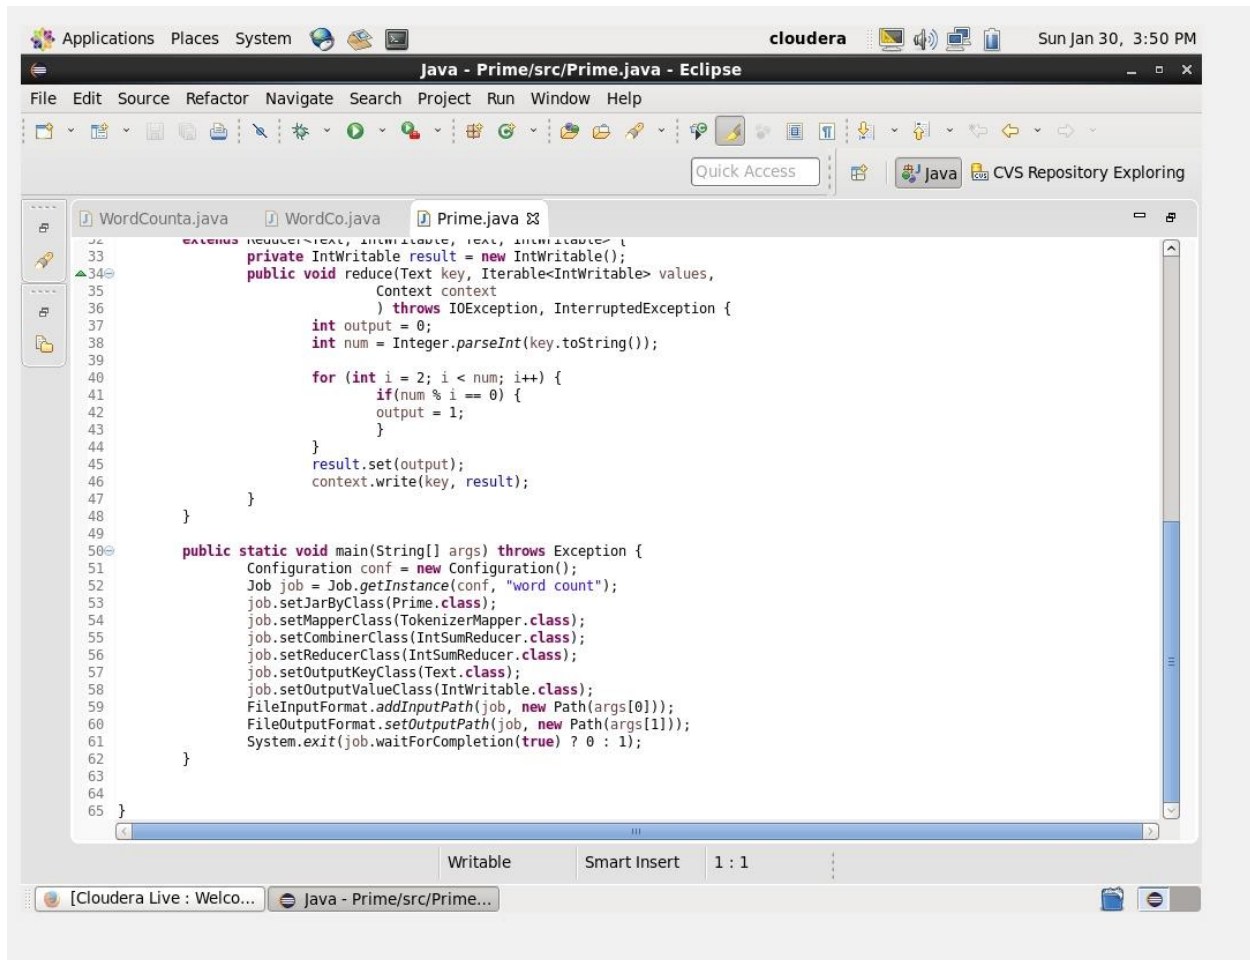


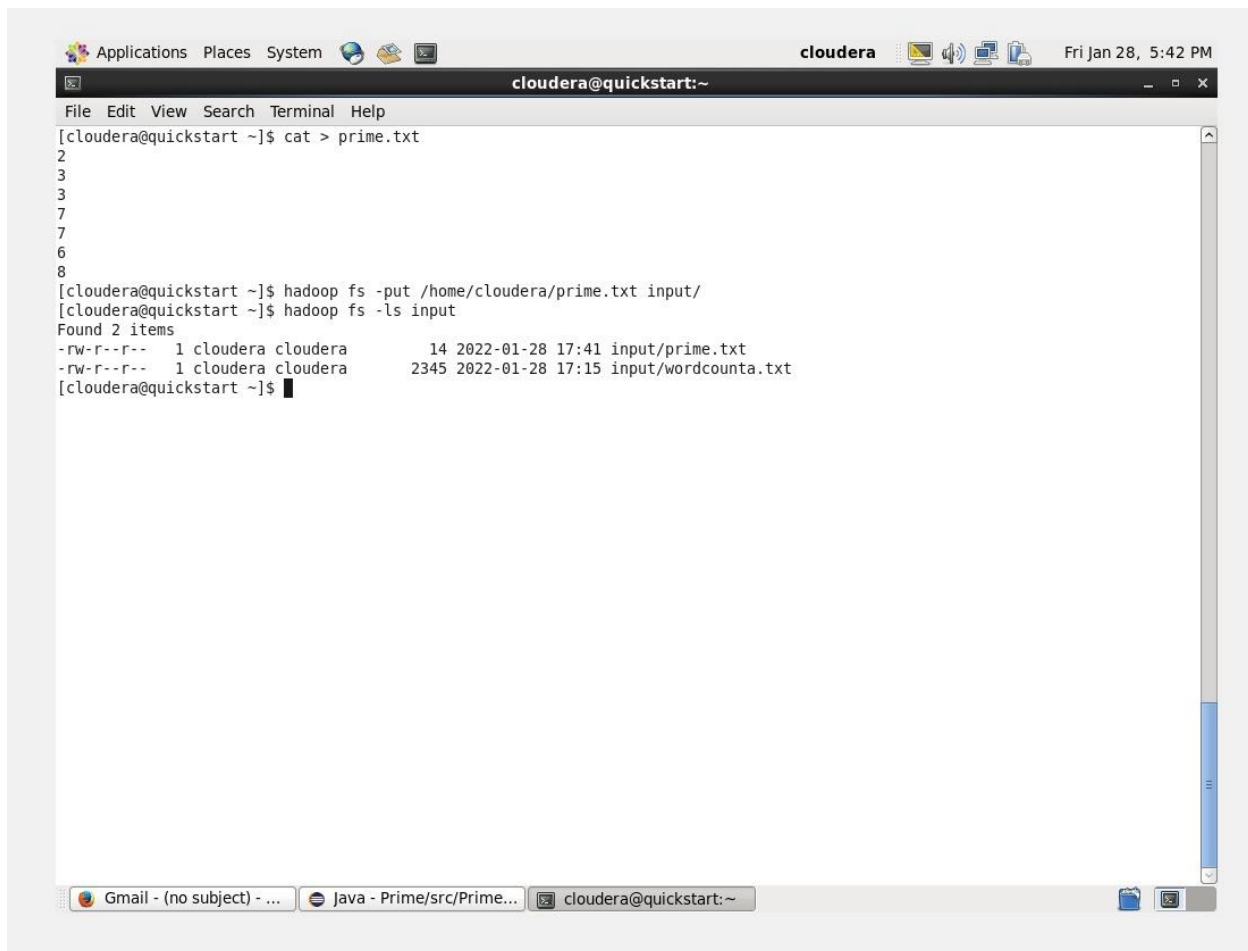
# Prime Number Program

```
1 import java.io.IOException;
2 import java.util.StringTokenizer;
3 import org.apache.hadoop.conf.Configuration;
4 import org.apache.hadoop.fs.Path;
5 import org.apache.hadoop.io.IntWritable;
6 import org.apache.hadoop.io.Text;
7 import org.apache.hadoop.mapreduce.Job;
8 import org.apache.hadoop.mapreduce.Mapper;
9 import org.apache.hadoop.mapreduce.Reducer;
10 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
11 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
12
13 public class Prime{
14
15     public static class TokenizerMapper
16         extends Mapper<Object, Text, Text, IntWritable> {
17
18         private final static IntWritable one = new IntWritable(1);
19         private Text word = new Text();
20
21     public void map(Object key, Text value, Context context
22         ) throws IOException, InterruptedException {
23         StringTokenizer itr = new StringTokenizer(value.toString());
24         while (itr.hasMoreTokens()) {
25             word.set(itr.nextToken());
26             context.write(word, one);
27         }
28     }
29
30     public static class IntSumReducer
31         extends Reducer<Text, IntWritable, Text, IntWritable> {
32         private IntWritable result = new IntWritable();
33         public void reduce(Text key, Iterable<IntWritable> values
```

- Firstly we create a java project named prime and then import all the required header files and add all external jar libs to the program.
- Then we extend the mapper class by creating TokenizerMapper and it takes arguments like key-in,value-in,key-out and value-out. Then we create word variable to store word and one variable to store value.
- Then we create map function which takes arguments like key-in as key, textual content as value variable and context to emit result to next stage.
- Then we define StringTokenizer which divides string into tokens and store each of them in itr variable.
- Then we define reducer class as IntSumReducer which takes four inputs. Then we create result variable to store how many number of times the key is appearing.
- Then we write logic which adds words which starts with letter b and gives output to the reducer.

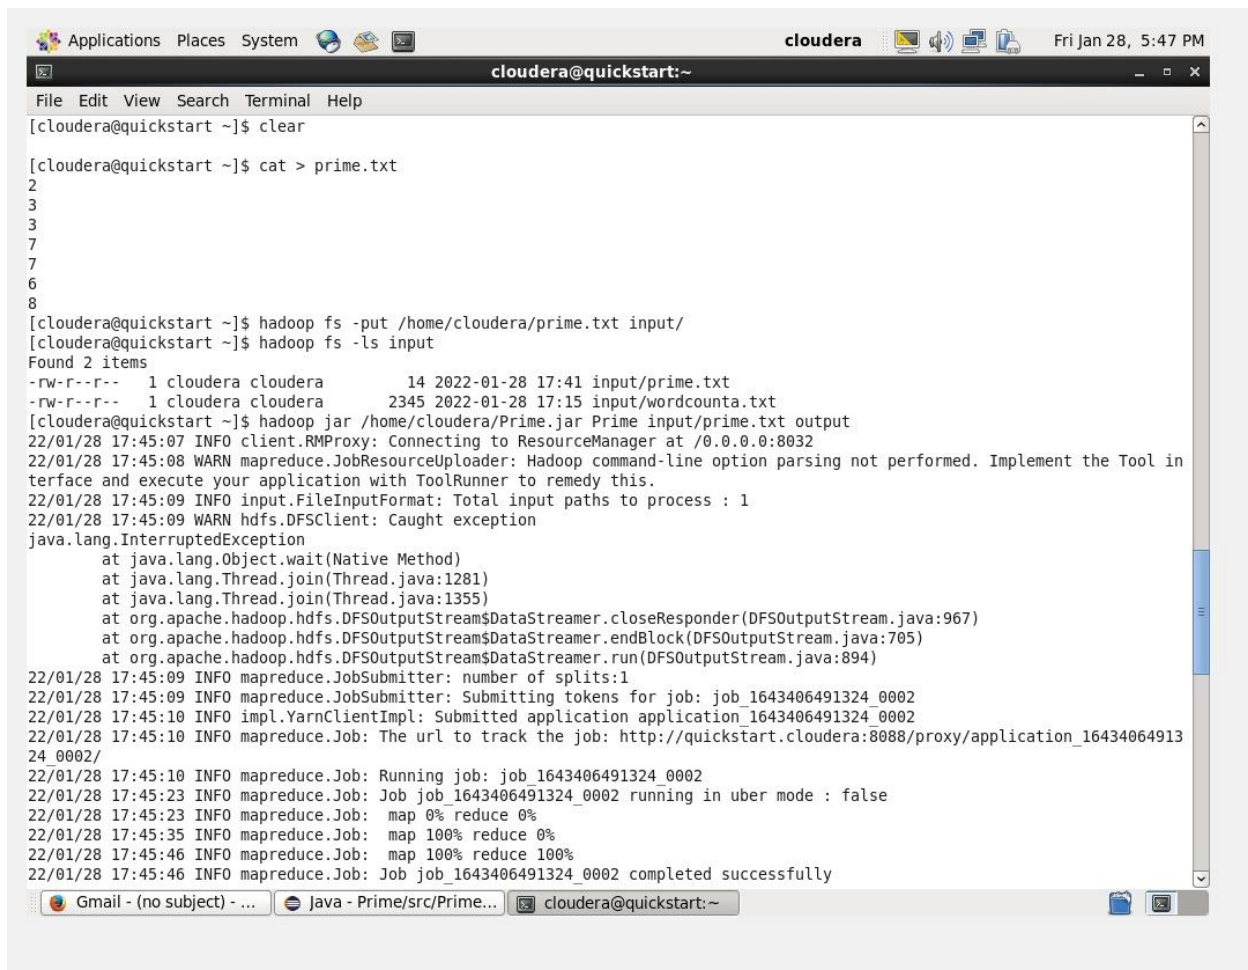






```
Applications Places System cloudera Fri Jan 28, 5:42 PM
cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ cat > prime.txt
2
3
3
7
7
6
8
[cloudera@quickstart ~]$ hadoop fs -put /home/cloudera/prime.txt input/
[cloudera@quickstart ~]$ hadoop fs -ls input
Found 2 items
-rw-r--r-- 1 cloudera cloudera      14 2022-01-28 17:41 input/prime.txt
-rw-r--r-- 1 cloudera cloudera    2345 2022-01-28 17:15 input/wordcounta.txt
[cloudera@quickstart ~]$
```

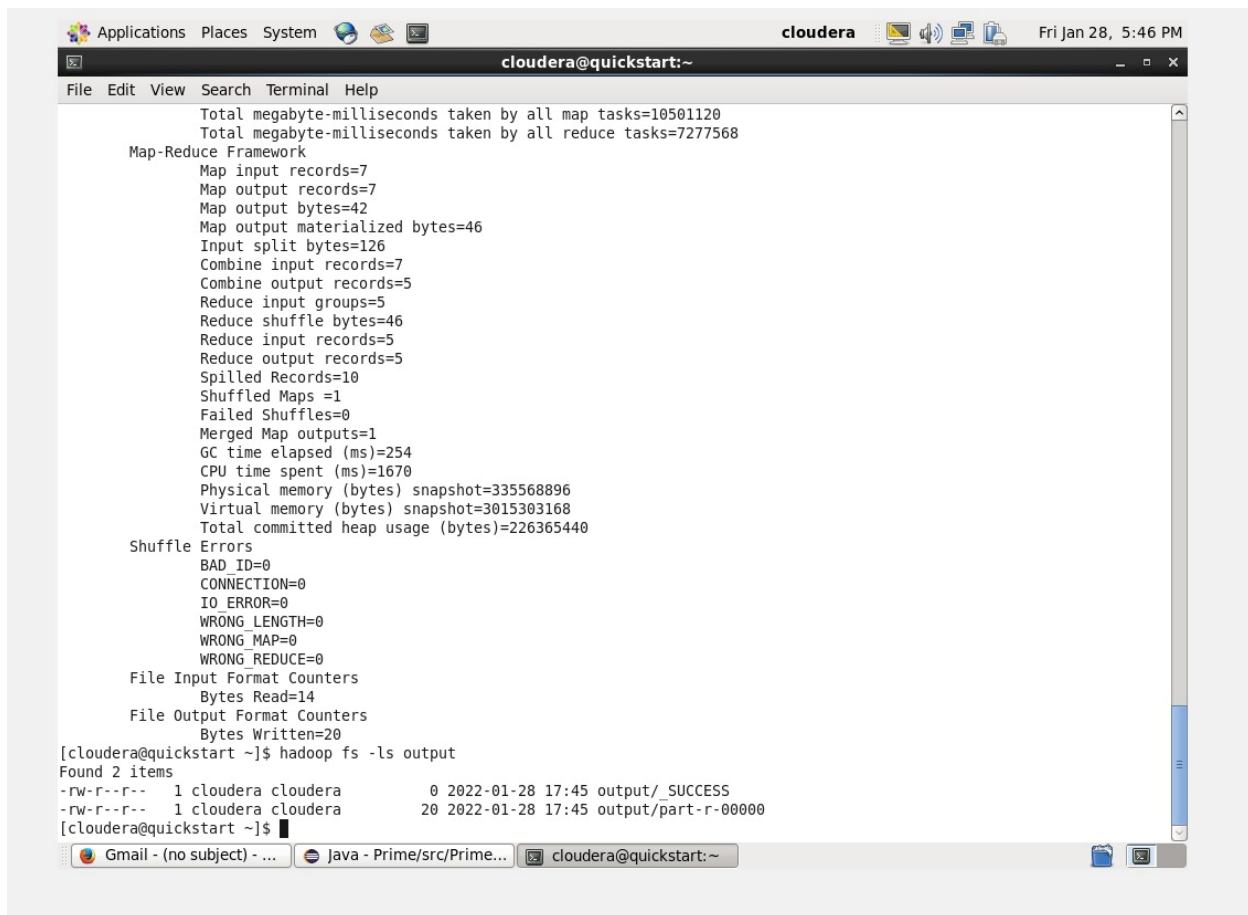
- Then I have added input to file using cat command and listed the files using ls command.



```
[cloudera@quickstart ~]$ clear

[cloudera@quickstart ~]$ cat > prime.txt
2
3
3
7
7
6
8
[cloudera@quickstart ~]$ hadoop fs -put /home/cloudera/prime.txt input/
[cloudera@quickstart ~]$ hadoop fs -ls input
Found 2 items
-rw-r--r-- 1 cloudera cloudera      14 2022-01-28 17:41 input/prime.txt
-rw-r--r-- 1 cloudera cloudera    2345 2022-01-28 17:15 input/wordcounta.txt
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Prime.jar Prime input/prime.txt output
22/01/28 17:45:07 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
22/01/28 17:45:08 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool in
terface and execute your application with ToolRunner to remedy this.
22/01/28 17:45:09 INFO input.FileInputFormat: Total input paths to process : 1
22/01/28 17:45:09 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/28 17:45:09 INFO mapreduce.JobSubmitter: number of splits:1
22/01/28 17:45:09 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1643406491324_0002
22/01/28 17:45:10 INFO impl.YarnClientImpl: Submitted application application_1643406491324_0002
22/01/28 17:45:10 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_16434064913
24_0002/
22/01/28 17:45:10 INFO mapreduce.Job: Running job: job_1643406491324_0002
22/01/28 17:45:23 INFO mapreduce.Job: Job job_1643406491324_0002 running in uber mode : false
22/01/28 17:45:23 INFO mapreduce.Job:  map 0% reduce 0%
22/01/28 17:45:35 INFO mapreduce.Job:  map 100% reduce 0%
22/01/28 17:45:46 INFO mapreduce.Job:  map 100% reduce 100%
22/01/28 17:45:46 INFO mapreduce.Job: Job job_1643406491324_0002 completed successfully
```

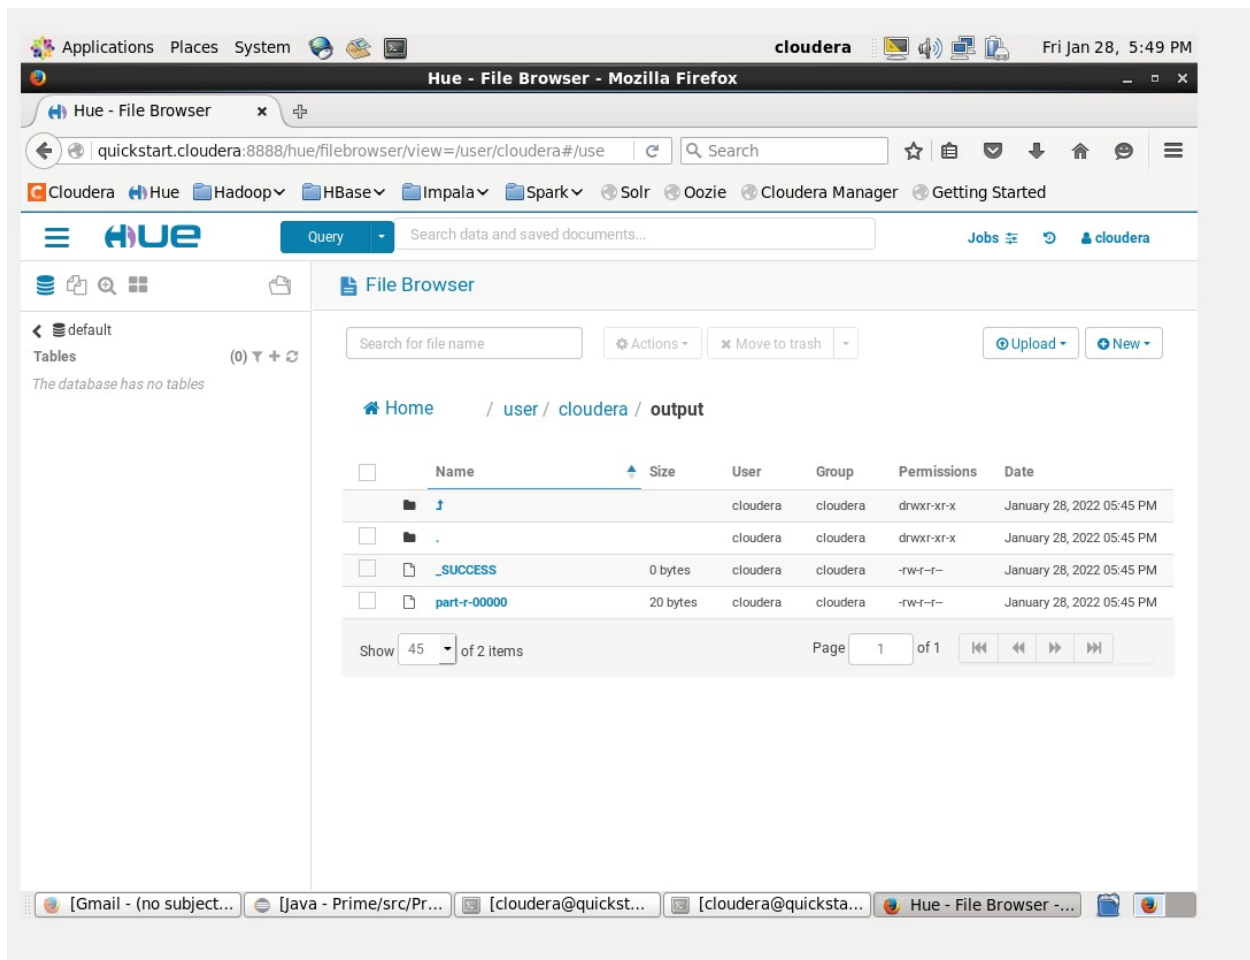
- Then I have run mapreduce job and executed the code.

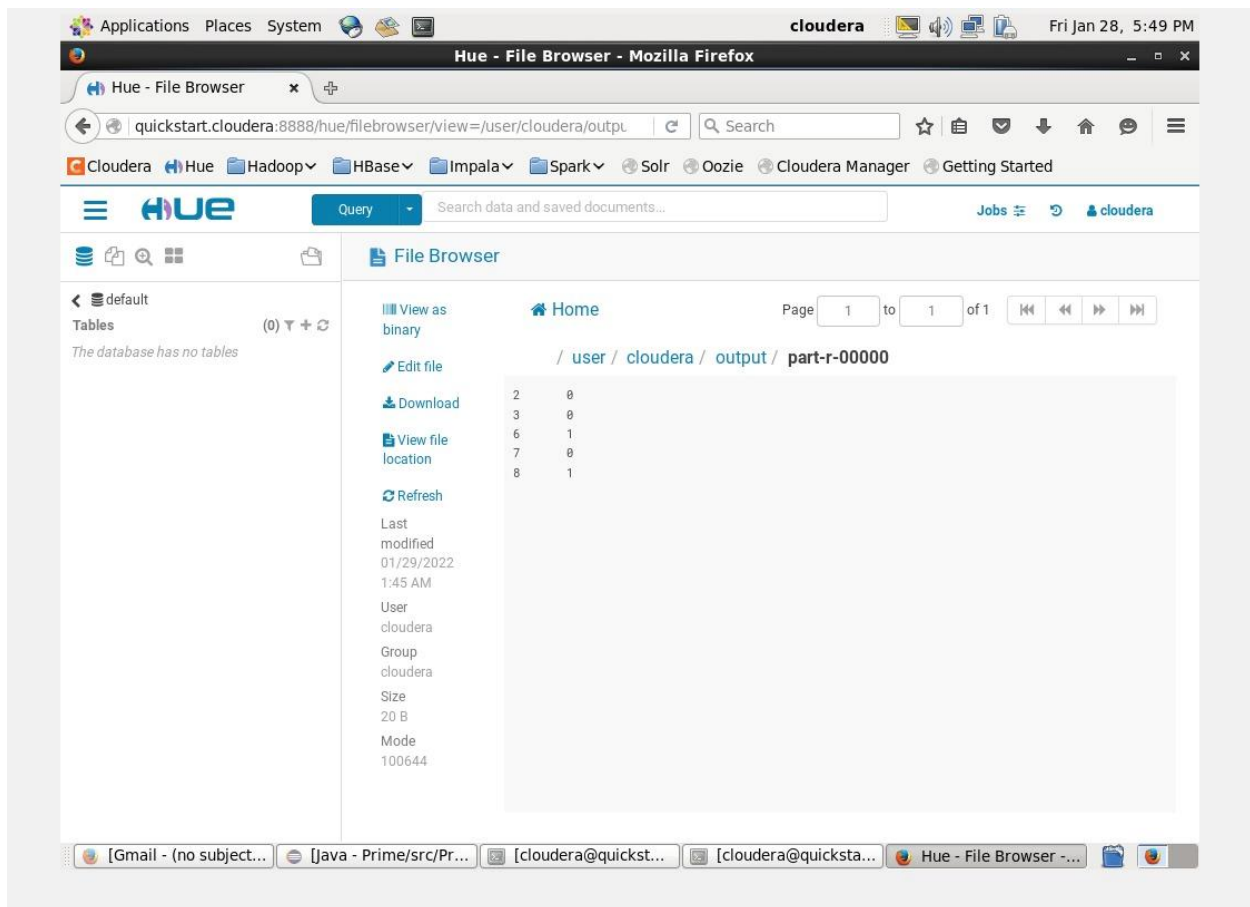


The screenshot shows a terminal window titled "cloudera@quickstart:~" with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal output displays the results of a Hadoop Map-Reduce job. It includes statistics for map and reduce tasks, a detailed breakdown of the Map-Reduce Framework (input/output records, bytes, shuffle, and reduce phases), shuffle errors (all zero), and file input/output format counters. Finally, it shows the output of the command "hadoop fs -ls output", which lists two files: "output/" and "output/part-r-00000".

```
cloudera@quickstart:~$ cat /dev/null
Total megabyte-milliseconds taken by all map tasks=10501120
Total megabyte-milliseconds taken by all reduce tasks=7277568
Map-Reduce Framework
  Map input records=7
  Map output records=7
  Map output bytes=42
  Map output materialized bytes=46
  Input split bytes=126
  Combine input records=7
  Combine output records=5
  Reduce input groups=5
  Reduce shuffle bytes=46
  Reduce input records=5
  Reduce output records=5
  Spilled Records=10
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=254
  CPU time spent (ms)=1670
  Physical memory (bytes) snapshot=335568896
  Virtual memory (bytes) snapshot=3015303168
  Total committed heap usage (bytes)=226365440
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=14
File Output Format Counters
  Bytes Written=20
[cloudera@quickstart ~]$ hadoop fs -ls output
Found 2 items
-rw-r--r--  1 cloudera cloudera      0 2022-01-28 17:45 output/ SUCCESS
-rw-r--r--  1 cloudera cloudera    20 2022-01-28 17:45 output/part-r-00000
[cloudera@quickstart ~]$
```

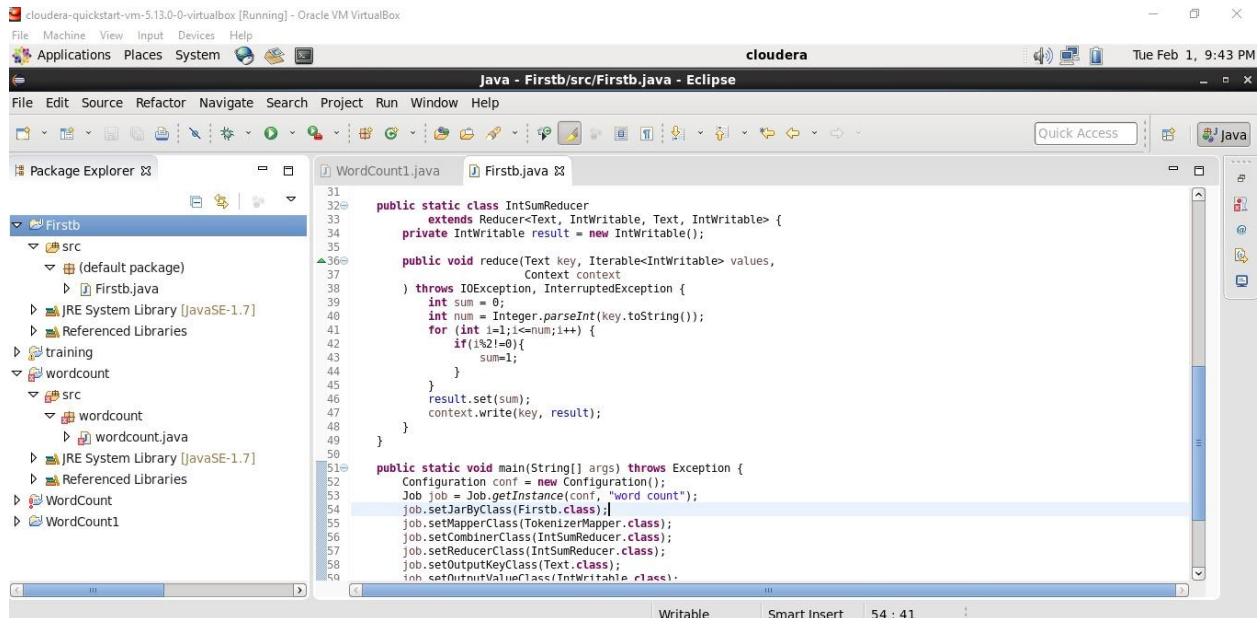
Then I displayed the output using ls command and it gives prime numbers as 1 and others as 0.





Here is visualized version of prime number program

# Odd number Program



```
31
32 public static class IntSumReducer
33     extends Reducer<Text, IntWritable, Text, IntWritable> {
34     private IntWritable result = new IntWritable();
35
36     public void reduce(Text key, Iterable<IntWritable> values,
37         Context context
38     ) throws IOException, InterruptedException {
39         int sum = 0;
40         int num = Integer.parseInt(key.toString());
41         for (int i=1; i<=num; i++) {
42             if (i%2!=0) {
43                 sum+=1;
44             }
45         }
46         result.set(sum);
47         context.write(key, result);
48     }
49 }
50
51 public static void main(String[] args) throws Exception {
52     Configuration conf = new Configuration();
53     Job job = Job.getInstance(conf, "word count");
54     job.setJarByClass(Firstb.class);
55     job.setMapperClass(TokenMapper.class);
56     job.setCombinerClass(IntSumReducer.class);
57     job.setReducerClass(IntSumReducer.class);
58     job.setOutputKeyClass(Text.class);
59     job.setOutputValueClass(IntWritable.class);
60 }
```

```
File Output Format Counters
  Bytes Written=12
[cloudera@quickstart ~]$ hadoop fs -ls out
Found 2 items
-rw-r--r--  1 cloudera cloudera      0 2022-02-01 21:27 out/_SUCCESS
-rw-r--r--  1 cloudera cloudera    12 2022-02-01 21:27 out/part-r-00000
[cloudera@quickstart ~]$ hadoop fs -cat out/part-r-00000
1      1
3      1
5      1
[cloudera@quickstart ~]$
```

